

STATS 500 HW8

Minxuan Chen

2023-11-16

Table of contents

Problem 1	1
Linear regression with all predictors	1
Linear regression with variables selected using AIC	3
Principal component regression - using CV to pick order of model	5
Partial least squares - using CV to pick order of model	7
Ridge regression - using GCV to pick regularization parameter	9
Lasso regression - using CV to pick regularization parameter t	12
Summary	13
Prediction	13

Github repo: https://github.com/PKUiiiiice/STATS_500

Problem 1

```
1 library(pls)
2 library(MASS)
3 library(glmnet)
4 library(lars)
5 gasdata = data.frame(cbind(gasoline$octane,
6                             gasoline$NIR[,c(1:40)]))
7 names(gasdata)=c('octane','NIR1','NIR2','NIR3','NIR4','NIR5','NIR6','NIR7','NIR8','NIR9',
8 'NIR10','NIR11','NIR12','NIR13','NIR14','NIR15','NIR16','NIR17','NIR18',
9 'NIR19','NIR20','NIR21','NIR22','NIR23','NIR24','NIR25','NIR26','NIR27',
10 'NIR28','NIR29','NIR30','NIR31','NIR32','NIR33','NIR34','NIR35','NIR36',
11 'NIR37','NIR38','NIR39','NIR40')
12 c = seq(1,56,5)
13 gasdata_tr = gasdata[-c,]
14 gasdata_te = gasdata[c,]
15
16 rmse <- function(x, y){
17   return(sqrt(mean((x-y)^2)))
18 }
```

Linear regression with all predictors

```
1 m.linear <- lm(octane ~ . , data=gasdata_tr)
2 summary(m.linear)
```

Call:

```
lm(formula = octane ~ . , data = gasdata_tr)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.43436	-0.13133	-0.04657	0.08413	0.44421

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	107.02	20.01	5.349	0.00107	**
NIR1	35.64	595.87	0.060	0.95398	
NIR2	384.21	1005.27	0.382	0.71365	
NIR3	10.17	1197.86	0.008	0.99346	
NIR4	-297.77	894.68	-0.333	0.74901	
NIR5	-93.43	710.43	-0.132	0.89907	
NIR6	-14.22	807.62	-0.018	0.98645	
NIR7	-35.36	587.05	-0.060	0.95365	
NIR8	-47.13	696.38	-0.068	0.94793	

NIR9	219.70	733.58	0.299	0.77326
NIR10	636.73	640.08	0.995	0.35299
NIR11	-354.73	850.22	-0.417	0.68901
NIR12	-905.76	958.04	-0.945	0.37593
NIR13	253.90	1035.57	0.245	0.81335
NIR14	-300.37	769.91	-0.390	0.70803
NIR15	-444.73	1209.94	-0.368	0.72406
NIR16	102.40	681.54	0.150	0.88481
NIR17	330.27	777.94	0.425	0.68391
NIR18	-986.47	1004.00	-0.983	0.35856
NIR19	-1102.90	1753.06	-0.629	0.54924
NIR20	-186.82	1692.29	-0.110	0.91519
NIR21	1767.13	1663.56	1.062	0.32339
NIR22	830.11	1632.29	0.509	0.62670
NIR23	558.40	1523.42	0.367	0.72479
NIR24	33.31	1073.68	0.031	0.97612
NIR25	-71.80	1115.86	-0.064	0.95050
NIR26	-415.67	1705.07	-0.244	0.81439
NIR27	1044.02	1327.17	0.787	0.45729
NIR28	-1774.28	2249.69	-0.789	0.45618
NIR29	1152.78	1935.89	0.595	0.57027
NIR30	-844.72	1964.14	-0.430	0.68007
NIR31	-701.52	2374.26	-0.295	0.77621
NIR32	473.06	2108.36	0.224	0.82888
NIR33	719.48	2467.20	0.292	0.77903
NIR34	1421.30	2230.36	0.637	0.54423
NIR35	-884.61	1775.08	-0.498	0.63351
NIR36	-1848.97	2428.44	-0.761	0.47130
NIR37	899.40	2267.59	0.397	0.70345
NIR38	304.03	1339.73	0.227	0.82696
NIR39	893.17	1379.99	0.647	0.53812
NIR40	-745.12	1682.32	-0.443	0.67119

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5232 on 7 degrees of freedom

Multiple R-squared: 0.983, Adjusted R-squared: 0.8858

F-statistic: 10.11 on 40 and 7 DF, p-value: 0.001893

```

1 #training error
2 paste("Training error",
3       tr.linear <- rmse(m.linear$fitted.values,
4                         gasdata_tr$octane))

```

```
[1] "Training error 0.199794670776533"
```

```

1 #test error
2 paste("Test error", te.linear <- rmse(predict(m.linear, newdata = gasdata_te),gasdata_te)

```

```
[1] "Test error 0.510418176467638"
```

Linear regression with variables selected using AIC

```

1 #AIC
2 step(m.linear, trace=FALSE)

```

Call:

```
lm(formula = octane ~ NIR2 + NIR4 + NIR10 + NIR12 + NIR14 + NIR15 +
  NIR17 + NIR18 + NIR19 + NIR21 + NIR23 + NIR27 + NIR28 + NIR29 +
  NIR30 + NIR31 + NIR33 + NIR34 + NIR36 + NIR38 + NIR39 + NIR40,
  data = gasdata_tr)
```

Coefficients:

(Intercept)	NIR2	NIR4	NIR10	NIR12	NIR14
99.7	684.4	-508.3	669.8	-744.8	-275.1
NIR15	NIR17	NIR18	NIR19	NIR21	NIR23
-484.8	370.3	-1117.2	-691.3	1460.4	595.2
NIR27	NIR28	NIR29	NIR30	NIR31	NIR33
965.0	-1330.5	830.0	-954.0	-1133.2	1231.0
NIR34	NIR36	NIR38	NIR39	NIR40	
1341.4	-1169.6	666.7	718.8	-1123.0	

```

1 m.AIC <- lm(octane ~ NIR2 + NIR4 + NIR10 + NIR12 +
2             NIR14 + NIR15 + NIR17 + NIR18 + NIR19 +
3             NIR21 + NIR23 + NIR27 + NIR28 + NIR29 +
4             NIR30 + NIR31 + NIR33 + NIR34 + NIR36 +
5             NIR38 + NIR39 + NIR40,
6             data = gasdata_tr)
7 summary(m.AIC)

```

Call:

```
lm(formula = octane ~ NIR2 + NIR4 + NIR10 + NIR12 + NIR14 + NIR15 +
  NIR17 + NIR18 + NIR19 + NIR21 + NIR23 + NIR27 + NIR28 + NIR29 +
  NIR30 + NIR31 + NIR33 + NIR34 + NIR36 + NIR38 + NIR39 + NIR40,
  data = gasdata_tr)
```

Residuals:

Min	1Q	Median	3Q	Max
-----	----	--------	----	-----

-0.42681 -0.14459 -0.05653 0.15007 0.51993

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	99.696	4.179	23.858	< 2e-16	***
NIR2	684.419	179.710	3.808	0.000809	***
NIR4	-508.335	104.534	-4.863	5.32e-05	***
NIR10	669.787	139.875	4.788	6.45e-05	***
NIR12	-744.776	266.674	-2.793	0.009874	**
NIR14	-275.066	252.737	-1.088	0.286823	
NIR15	-484.831	335.843	-1.444	0.161260	
NIR17	370.275	281.261	1.316	0.199955	
NIR18	-1117.209	319.204	-3.500	0.001766	**
NIR19	-691.291	356.085	-1.941	0.063567	.
NIR21	1460.441	456.963	3.196	0.003753	**
NIR23	595.236	406.480	1.464	0.155556	
NIR27	965.036	334.688	2.883	0.007976	**
NIR28	-1330.519	533.443	-2.494	0.019596	*
NIR29	830.047	482.421	1.721	0.097682	.
NIR30	-954.016	622.262	-1.533	0.137800	
NIR31	-1133.170	568.326	-1.994	0.057183	.
NIR33	1230.981	546.001	2.255	0.033172	*
NIR34	1341.382	503.985	2.662	0.013396	*
NIR36	-1169.612	659.330	-1.774	0.088255	.
NIR38	666.669	380.654	1.751	0.092142	.
NIR39	718.798	433.285	1.659	0.109621	
NIR40	-1123.047	594.768	-1.888	0.070656	.

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3035 on 25 degrees of freedom

Multiple R-squared: 0.9796, Adjusted R-squared: 0.9616

F-statistic: 54.44 on 22 and 25 DF, p-value: 5.448e-16

```
1 #training error
2 paste("Training error", tr.AIC <- rmse(m.AIC$fitted.values,
3 gasdata_tr$octane))
```

```
[1] "Training error 0.219023876211711"
```

```
1 #test error
2 paste("Test error", te.AIC <- rmse(predict(m.AIC, newdata = gasdata_te),
3 gasdata_te$octane))
```

```
[1] "Test error 0.665773147711337"
```

Principal component regression - using CV to pick order of model

```

1 #PCR
2 m.PCR <- pcr(octane~., data=gasdata_tr, validation = "CV")
3 summary(m.PCR)

```

Data: X dimension: 48 40
Y dimension: 48 1
Fit method: svdpc
Number of components considered: 40

VALIDATION: RMSEP

Cross-validated using 10 random segments.

	(Intercept)	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps
CV	1.564	1.581	1.135	0.8511	1.055	1.353	1.502
adjCV	1.564	1.578	1.068	0.8450	1.037	1.319	1.461
	7 comps	8 comps	9 comps	10 comps	11 comps	12 comps	13 comps
CV	1.493	1.487	1.547	1.476	1.423	1.428	1.426
adjCV	1.446	1.443	1.501	1.416	1.370	1.376	1.373
	14 comps	15 comps	16 comps	17 comps	18 comps	19 comps	20 comps
CV	1.363	1.359	1.363	1.411	1.434	1.438	1.415
adjCV	1.309	1.305	1.309	1.355	1.378	1.381	1.359
	21 comps	22 comps	23 comps	24 comps	25 comps	26 comps	27 comps
CV	1.361	1.354	1.339	1.305	1.305	1.272	1.269
adjCV	1.305	1.298	1.284	1.251	1.252	1.219	1.217
	28 comps	29 comps	30 comps	31 comps	32 comps	33 comps	34 comps
CV	1.274	1.307	1.404	1.376	1.398	1.478	1.516
adjCV	1.222	1.255	1.348	1.319	1.340	1.417	1.452
	35 comps	36 comps	37 comps	38 comps	39 comps	40 comps	
CV	1.636	1.688	1.966	2.128	2.193	3.085	
adjCV	1.564	1.614	1.875	2.029	2.091	2.936	

TRAINING: % variance explained


	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps
X	96.6413	98.08	99.31	99.64	99.72	99.79	99.84	99.88
octane	0.1102	61.00	72.79	81.89	82.05	83.40	85.88	86.21
	9 comps	10 comps	11 comps	12 comps	13 comps	14 comps	15 comps	
X	99.90	99.92	99.93	99.94	99.95	99.96	99.96	
octane	86.62	94.45	94.47	94.47	94.99	95.85	96.03	
	16 comps	17 comps	18 comps	19 comps	20 comps	21 comps	22 comps	
X	99.97	99.97	99.98	99.98	99.98	99.99	99.99	
octane	96.04	96.08	96.20	96.23	96.68	97.08	97.09	
	23 comps	24 comps	25 comps	26 comps	27 comps	28 comps	29 comps	
X	99.99	99.99	99.99	99.99	100.00	100.00	100.00	
octane	97.10	97.26	97.26	97.35	97.37	97.39	97.39	
	30 comps	31 comps	32 comps	33 comps	34 comps	35 comps	36 comps	

X	100.0	100.00	100.00	100.00	100.00	100.00	100.00
octane	97.4	97.84	97.87	97.88	97.88	98.06	98.06
	37 comps	38 comps	39 comps	40 comps			
X	100.00	100.0	100.00	100.0			
octane	98.14	98.2	98.27	98.3			

```
1 rmsCV <- RMSEP(m.PCR, estimate='CV')
2 which.min(rmsCV$val)
```

[1] 4

```
1 #plot
2 plot(rmsCV$val, xlab="PC number", ylab="CV RMS")
```



```
1 m.PCR.best <- pcr(octane~., data=gasdata_tr, ncomp=4, validation = "CV")
2 summary(m.PCR.best)
```

Data: X dimension: 48 40

Y dimension: 48 1

Fit method: svdpc

Number of components considered: 4

VALIDATION: RMSEP

Cross-validated using 10 random segments.

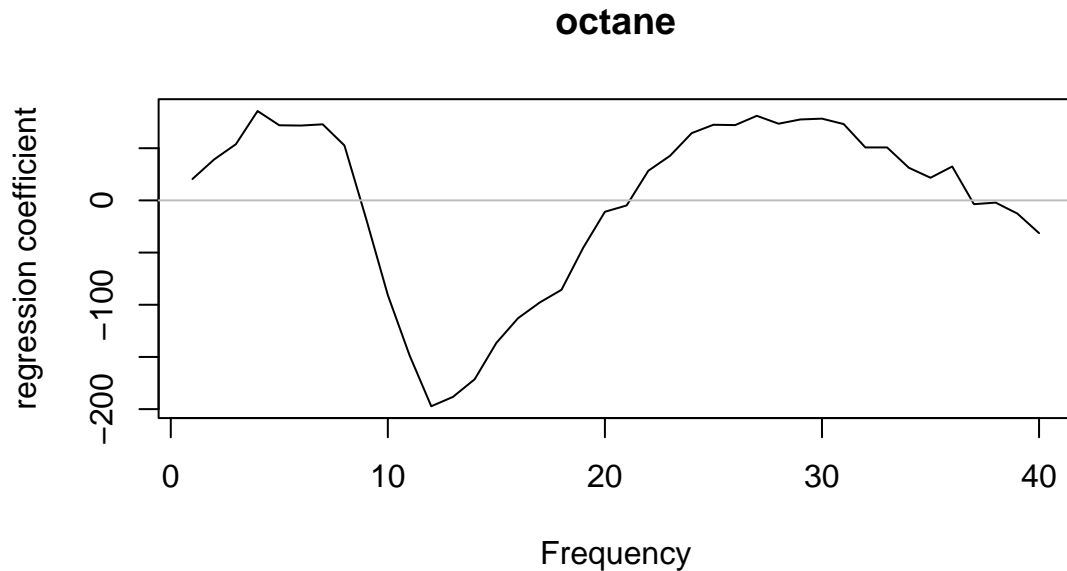
	(Intercept)	1 comps	2 comps	3 comps	4 comps
CV	1.564	1.584	1.05	0.9144	1.161
adjCV	1.564	1.581	1.02	0.9055	1.132

TRAINING: % variance explained

1 comps	2 comps	3 comps	4 comps
---------	---------	---------	---------

X	96.6413	98.08	99.31	99.64
octane	0.1102	61.00	72.79	81.89

```
1 coefplot(m.PCR.best, ncomp=4, xlab="Frequency")
```



```
1 pred.pctr <- predict(m.PCR, newdata=gasdata_tr, ncomp=4)
2 pred.pcte <- predict(m.PCR, newdata=gasdata_te, ncomp=4)
3 #training error
4 paste("Training error", tr.PCR <- rmse(pred.pctr,
5                                         gasdata_tr$octane))
```

```
[1] "Training error 0.651816057353665"
```

```
1 #test error
2 paste("Test error", te.PCR <- rmse(pred.pcte, gasdata_te$octane))
```

```
[1] "Test error 0.421555414203408"
```

Partial least squares - using CV to pick order of model

```
1 m.pls <- pls(octane ~ ., data = gasdata_tr,
2             validation = "CV")
3 summary(m.pls)
```


Data: X dimension: 48 40
Y dimension: 48 1
Fit method: kernelpls
Number of components considered: 40

VALIDATION: RMSEP

Cross-validated using 10 random segments.

	(Intercept)	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps
CV	1.564	1.691	0.8789	0.9622	1.427	1.418	1.390
adjCV	1.564	1.636	0.8717	0.9476	1.377	1.357	1.324
	7 comps	8 comps	9 comps	10 comps	11 comps	12 comps	13 comps
CV	1.366	1.355	1.361	1.44	1.490	1.506	1.565
adjCV	1.301	1.291	1.295	1.37	1.417	1.432	1.488
	14 comps	15 comps	16 comps	17 comps	18 comps	19 comps	20 comps
CV	1.638	1.717	1.804	1.856	1.879	1.935	1.981
adjCV	1.557	1.633	1.714	1.764	1.786	1.839	1.884
	21 comps	22 comps	23 comps	24 comps	25 comps	26 comps	27 comps
CV	1.994	2.055	2.108	2.144	2.214	2.253	2.319
adjCV	1.896	1.955	2.006	2.040	2.106	2.143	2.205
	28 comps	29 comps	30 comps	31 comps	32 comps	33 comps	34 comps
CV	2.385	2.437	2.517	2.625	2.739	2.815	2.923
adjCV	2.268	2.316	2.392	2.494	2.602	2.673	2.775
	35 comps	36 comps	37 comps	38 comps	39 comps	40 comps	
CV	3.007	3.042	3.074	3.094	3.103	3.150	
adjCV	2.854	2.887	2.917	2.937	2.945	2.989	


TRAINING: % variance explained

	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps
X	84.87	98.05	98.72	99.63	99.69	99.73	99.78	99.84
octane	11.50	76.97	83.03	85.21	92.64	95.68	96.49	96.73
	9 comps	10 comps	11 comps	12 comps	13 comps	14 comps	15 comps	
X	99.86	99.88	99.90	99.93	99.93	99.94	99.95	
octane	97.24	97.45	97.64	97.74	97.90	98.01	98.06	
	16 comps	17 comps	18 comps	19 comps	20 comps	21 comps	22 comps	
X	99.95	99.96	99.96	99.96	99.97	99.98	99.98	
octane	98.10	98.14	98.17	98.21	98.23	98.25	98.26	
	23 comps	24 comps	25 comps	26 comps	27 comps	28 comps	29 comps	
X	99.98	99.98	99.99	99.99	99.99	99.99	99.99	
octane	98.26	98.27	98.29	98.29	98.29	98.30	98.30	
	30 comps	31 comps	32 comps	33 comps	34 comps	35 comps	36 comps	
X	99.99	99.99	100.0	100.0	100.0	100.0	100.0	
octane	98.30	98.30	98.3	98.3	98.3	98.3	98.3	
	37 comps	38 comps	39 comps	40 comps				
X	100.0	100.0	100.0	100.0				
octane	98.3	98.3	98.3	98.3				

```

1 plrmscv <- RMSEP(m.pls,estimate='CV')
2 plot(plrmscv$val)

```



```

1 which.min(plrmscv$val)

```

```
[1] 3
```

```

1 pred.plstr = predict(m.pls, newdata=gasdata_tr, ncomp=3)
2 pred.plste = predict(m.pls, newdata=gasdata_te, ncomp=3)
3
4 #training error
5 paste("Training error", tr.pls <- rmse(pred.plstr,
6                                         gasdata_tr$octane))

```

```
[1] "Training error 0.6309353614644"
```

```

1 #test error
2 paste("Test error", te.pls <- rmse(pred.plste,gasdata_te$octane))

```

```
[1] "Test error 0.32577507300782"
```

Ridge regression - using GCV to pick regularization parameter

```

1 m.ridge <- lm.ridge(octane ~ .,
2                     lambda=seq(0, .2, .01),
3                     data = gasdata_tr)
4 head(m.ridge$coef)


```

	0.00	0.01	0.02	0.03	0.04	0.05
NIR1	0.16119223	-0.4436893	-0.57325342	-0.6232798	-0.6451038	-0.6533857
NIR2	1.70484815	1.7086173	1.56902266	1.4791696	1.4128115	1.3603658
NIR3	0.04577975	0.7002692	0.81807474	0.8413829	0.8367568	0.8215341
NIR4	-1.42156175	-1.9007774	-1.63293466	-1.4617888	-1.3408500	-1.2492333
NIR5	-0.43953606	0.1809873	-0.02078692	-0.1427672	-0.2191168	-0.2686981
NIR6	-0.06985119	-0.9974199	-1.08400188	-1.0876496	-1.0697388	-1.0449144
	0.06	0.07	0.08	0.09	0.10	0.11
NIR1	-0.6540934	-0.6501841	-0.6433181	-0.6345095	-0.6244151	-0.6134783
NIR2	1.3170269	1.2800392	1.2476937	1.2188703	1.1928032	1.1689502
NIR3	0.8020527	0.7810546	0.7598394	0.7390505	0.7190060	0.6998534
NIR4	-1.1764229	-1.1164996	-1.0658482	-1.0221226	-0.9837278	-0.9495380
NIR5	-0.3014929	-0.3231959	-0.3372779	-0.3459666	-0.3507517	-0.3526627
NIR6	-1.0180082	-0.9908823	-0.9643080	-0.9386054	-0.9138929	-0.8901947
	0.12	0.13	0.14	0.15	0.16	0.17
NIR1	-0.6020081	-0.5902244	-0.5782861	-0.5663094	-0.5543795	-0.5425594
NIR2	1.1469160	1.1264051	1.1071916	1.0890992	1.0719879	1.0557448
NIR3	0.6816481	0.6643943	0.6480682	0.6326309	0.6180362	0.6042352
NIR4	-0.9187347	-0.8907071	-0.8649900	-0.8412221	-0.8191184	-0.7984508
NIR5	-0.3524303	-0.3505844	-0.3475168	-0.3435214	-0.3388218	-0.3335899
NIR6	-0.8674908	-0.8457405	-0.8248946	-0.8049016	-0.7857107	-0.7672733
	0.18	0.19	0.20			
NIR1	-0.5308952	-0.5194203	-0.5081589			
NIR2	1.0402773	1.0255086	1.0113742			
NIR3	0.5911786	0.5788184	0.5671088			
NIR4	-0.7790344	-0.7607174	-0.7433744			
NIR5	-0.3279592	-0.3220339	-0.3158964			
NIR6	-0.7495437	-0.7324793	-0.7160407			

```

1  matplot(m.ridge$lambda,
2         t(m.ridge$coef), type="l",
3         lty=1,xlab=expression(lambda),
4         ylab=expression(hat(beta)))

```



```

1  ## Select an appropriate lambda
2  select(m.ridge)

```

modified HKB estimator is 0.02736696
modified L-W estimator is 4.510166
smallest value of GCV at 0.07

```

1  m.ridge.best <- lm.ridge(octane ~ .,
2                          lambda=0.07,
3                          data = gasdata_tr); m.ridge.best$coef

```

NIR1	NIR2	NIR3	NIR4	NIR5	NIR6
-0.65018407	1.28003922	0.78105457	-1.11649959	-0.32319592	-0.99088228
NIR7	NIR8	NIR9	NIR10	NIR11	NIR12
0.77130035	0.70097931	1.12018974	1.19684236	-2.28998131	-1.11024805
NIR13	NIR14	NIR15	NIR16	NIR17	NIR18
-0.41630192	-1.07836555	-0.06183079	-1.83155778	0.45444323	-0.87501642
NIR19	NIR20	NIR21	NIR22	NIR23	NIR24
-0.47712229	-0.51357628	1.11889820	0.38226678	1.47916562	1.10589907
NIR25	NIR26	NIR27	NIR28	NIR29	NIR30
0.83503834	0.39766529	1.63225580	-1.12931271	-1.29815497	-1.19598157
NIR31	NIR32	NIR33	NIR34	NIR35	NIR36
-0.90783681	0.61638616	0.44403495	1.17419805	0.54167824	-0.29424581
NIR37	NIR38	NIR39	NIR40		
1.00761541	0.51718366	0.20540635	-1.36124868		

```

1  yfitttr <- m.ridge.best$ym + scale(gasdata_tr[, -1],
2                                   center=m.ridge.best$xm,
3                                   scale=m.ridge.best$scales) %*%
4                                   m.ridge.best$coef
5  # training error
6  paste("Training error", tr.ridge <- rmse(yfitttr,
7                                           gasdata_tr$octane))

```

```
[1] "Training error 0.279915048711154"
```

```

1  yfritte <- m.ridge.best$ym + scale(gasdata_te[, -1],
2                                   center=m.ridge.best$xm,
3                                   scale=m.ridge.best$scales) %*%
4                                   m.ridge.best$coef
5  #test error
6  paste("Test error", te.ridge <- rmse(yfritte,


```

```
7 gasdata_te$octane))
```

```
[1] "Test error 0.491558224890926"
```

Lasso regression - using CV to pick regularization parameter t.

```
1 set.seed(123)
2 m.lasso <- lars(as.matrix(gasdata_tr[,-1]),
3                 gasdata_tr$octane)
4 cvout <- cv.lars(as.matrix(gasdata_tr[,-1]),
5                  gasdata_tr$octane)
```



```
1 cvout$index[which.min(cvout$cv)]
```

```
[1] 0.05050505
```

```
1 m.lasso <- lars(as.matrix(gasdata_tr[,-1]),gasdata_tr$octane)
2 predlars_tr <- predict(m.lasso,as.matrix(gasdata_tr[,-1]),
3                        s=cvout$index[which.min(cvout$cv)],
4                        mode="fraction")$fit
5 # training error
6 paste("Training error", tr.lasso <- rmse(predlars_tr, gasdata_tr$octane))
```

```
[1] "Training error 0.826073321265821"
```

```
1 predlars_te <- predict(m.lasso,as.matrix(gasdata_te[,-1]),
2                        s=cvout$index[which.min(cvout$cv)],
```

```

3             mode="fraction")$fit
4 #test error
5 paste("Test error", te.lasso <- rmse(predlars_te, gasdata_te$octane))

```

```
[1] "Test error 0.646539408291542"
```

Summary

```

1 rmse.data <- matrix(c(tr.linear, tr.AIC, tr.PCR, tr.pls, tr.ridge, tr.lasso,
2                       te.linear, te.AIC, te.PCR, te.pls, te.ridge, te.lasso),
3                     ncol=6, byrow=1)
4 colnames(rmse.data) <- c("LR", "LR.AIC", "PCR",
5                           "PLS", "Ridge", "Lasso")
6 row.names(rmse.data) <- c("Training RMSE", "Test RMSE")
7
8 print(rmse.data)

```

	LR	LR.AIC	PCR	PLS	Ridge	Lasso
Training RMSE	0.1997947	0.2190239	0.6518161	0.6309354	0.2799150	0.8260733
Test RMSE	0.5104182	0.6657731	0.4215554	0.3257751	0.4915582	0.6465394

Prediction

```

1 meanfreq <- data.frame(t(colMeans(gasdata[,2:41])))
2
3 pred.l <- predict(m.linear, newdata=meanfreq)
4 pred.aic <- predict(m.AIC, newdata=meanfreq)
5 pred.pcr <- predict(m.PCR.best, newdata=meanfreq, ncomp=4)
6 pred.pls <- predict(m.pls, newdata=meanfreq, ncomp=3)
7 pred.ridge <- m.ridge.best$ym +
8   scale(meanfreq, center=m.ridge.best$xm,
9         scale=m.ridge.best$scales)%*%m.ridge.best$coef
10 pred.lasso <- predict(m.lasso, as.matrix(meanfreq), s=0.05050505, mode="fraction")$fit
11 c(pred.l, pred.aic, pred.pcr, pred.pls, pred.ridge, pred.lasso)

```

	1	1
	87.11089	87.13930
	87.13647	87.15794
	87.13247	87.16033